

dplyr Tutorial

Freshwater Ecology Research Group

Justin Pomeranz

March 30, 2017

This tutorial is aimed at giving you an introduction to the popular **dplyr** package. **dplyr** provides a few simple verbs that allow you to quickly and easily select and manipulate your data, and create an interactive environment for data exploration. This tutorial is based on the introduction and two-table vignettes for the **dplyr** package:

```
vignette("introduction", "dplyr")  
vignette("two-table", "dplyr")
```

As well as the book “*R for Data Science* by Hadley Wickham & Garrett Grolemund (O’reilly). Copyright 2017. 978-1-491-91039-9.” Available free online at <http://r4ds.had.co.nz/>

At the end, I will make a brief plug for the **tidyr** package, which is designed to make messy datasets into “tidy” data sets.

Both **dplyr** and **tidyr** are part of the suite of packages known as **tidyverse**. The packages within the tidyverse were designed with a similar data and programming philosophy, and work together fluidly. Hadley Wickham is the lead developer of the tidyverse, and his book above takes a tidyverse-centric view, and is a fantastic resource if you want to learn more.

Part 1: dplyr

Data manipulation is one of the primary tasks that scientists undertake. the **dplyr** package makes it easy to select the data you want, organize it in a useful way, and calculate useful new variables.

There is an extremely helpful Rstudio data wrangling cheat sheet available at <https://www.rstudio.com/wp-content/uploads/2015/02/data-wrangling-cheatsheet.pdf/>. I almost always have this cheat sheet open while I’m conducting any analyses. It takes a little effort to learn how to “read” the cheat sheet, but it’s well worth the effort.

First, make sure that the **dplyr** and **tidyr** packages are installed and loaded. We will also be using the **nycflights13** library as an example data set, and use **ggplot2** to illustrate an example.

```
install.packages(c("dplyr", "tidyr", "ggplot2", "nycflights13"))  
library(dplyr)  
library(tidyr)  
library(ggplot2)  
library(nycflights13)
```

Alternatively, you can install and load the **tidyverse** package, which includes **dplyr**, **tidyr**, **ggplot2** as well as many others. You will still need to install and load the **nycflights13** package separately.

```
install.packages("tidyverse")  
library(tidyverse)
```

Data: nycflights13

In order to demonstrate the basic data manipulation verbs of dplyr, we will be using the `nycflights13` database. This library contains three dataframes: `flights`, `airlines`, `airports`. The `flights` dataset contains all 336776 flights that departed from New York City in 2013. The data comes from the US Bureau of Transportation Statistics, and is documented in `?nycflights13::flights`

```
# make sure the nycflights13 library is loaded
library(nycflights13)
dim(flights)
#> [1] 336776      19
head(flights)
#> # A tibble: 6 × 19
#>   year month   day dep_time sched_dep_time dep_delay arr_time
#>   <int> <int> <int>   <int>         <int>      <dbl>   <int>
#> 1  2013     1     1     517             515         2     830
#> 2  2013     1     1     533             529         4     850
#> 3  2013     1     1     542             540         2     923
#> 4  2013     1     1     544             545        -1    1004
#> # ... with 2 more rows, and 12 more variables: sched_arr_time <int>,
#> #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
#> #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
#> #   minute <dbl>, time_hour <dtm>
```

You may notice that the `head(flights)` output is different from many other data frames that you may have worked with before. It only shows the rows and columns that easily fit into your window. That's because the `flights` object is a tibble:

```
class(flights)
#> [1] "tbl_df"      "tbl"        "data.frame"
```

Tibbles "tbl_df" act just as a data frame would, but they print only the information that can fit into your window. To see the full data frame in the Rstudio viewer use:

```
View(flights)
```

Or as a normal data frame use:

```
head(as.data.frame(flights))
```

To convert a data frame to a tibble, use `as_tibble()`.

Single table verbs

dplyr aims to provide a function for each basic verb of data manipulation:

- Pick observations based on their values: `filter()`
- Reorder the rows: `arrange()`
- Pick variables based on their names: `select()`
- Add new variables using functions on existing variables: `mutate()`
- Reduce many observations to a single summary: `summarise()`

All of these verbs can also be used with `group_by()`, which allows you to apply a verb group-by-group instead of on the whole dataset but more on this later.

All verbs work similarly and have similar syntaxes:

1. The first argument is a data frame.

2. Following arguments tell what to do with the data frame using un-quoted variable names.
3. Result is new data frame.

There are other verbs in the `dplyr` package, along with many useful “helper” functions. Browse the vignettes for `dplyr`: `browseVignettes(package = "dplyr")`, or check the Rstudio data wrangling cheat sheet <https://www.rstudio.com/wp-content/uploads/2015/02/data-wrangling-cheatsheet.pdf>

Filter rows with `filter()`

`filter()` allows you to select a subset of rows in a data frame. The first argument is the name of the data frame. The second and subsequent arguments are the expressions that filter the data frame:

For example, we can select all flights on July 4th with:

```
filter(flights, month == 7, day == 4)
#> # A tibble: 737 × 19
#>   year month   day dep_time sched_dep_time dep_delay arr_time
#>   <int> <int> <int>   <int>         <int>      <dbl>   <int>
#> 1  2013     7     4       11           2359         12     400
#> 2  2013     7     4        59           2359         60     501
#> 3  2013     7     4      454           500         -6     635
#> 4  2013     7     4      535           536         -1     802
#> # ... with 733 more rows, and 12 more variables: sched_arr_time <int>,
#> #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
#> #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
#> #   minute <dbl>, time_hour <dtm>
```

This is equivalent to the standard subsetting notation in base R:

```
flights[flights$month == 7 & flights$day == 4, ]
```

`filter()` works similarly to `subset()` but is slightly more flexible, where an `&` can be replaced with a `,:`

```
# flights between March and September
# ', '
filter(flights, month >= 3, month <= 9)
# '&'
filter(flights, month >= 3 & month <= 9)
```

You can also use other boolean operators, singly and in combination:

```
# flights in January or February
filter(flights, month == 1 | month == 2)
# flights in January or February that flew > 1000 or < 500 miles
filter(flights, month == 1 | month == 2 & distance > 1000 | distance < 500)
```

Arrange rows with `arrange()`

`arrange()` has a similar syntax as `filter()` but instead of filtering or selecting rows, it reorders them. It takes a data frame, and column name(s) to order by as arguments. The data will be arranged by the first column name provided, with ties being broken by subsequent columns:

```
# arrange by sched_dep_time
arrange(flights, sched_dep_time)
#> # A tibble: 336,776 × 19
#>   year month   day dep_time sched_dep_time dep_delay arr_time
```

```

#>   <int> <int> <int>      <int>          <int>      <dbl>      <int>
#> 1  2013     7    27        NA            106        NA        NA
#> 2  2013     1     2        458           500        -2        703
#> 3  2013     1     3        458           500        -2        650
#> 4  2013     1     4        456           500        -4        631
#> # ... with 336,772 more rows, and 12 more variables: sched_arr_time <int>,
#> #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
#> #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
#> #   minute <dbl>, time_hour <dtm>
# arrange by sched_dep_time, then dep_delay
arrange(flights, sched_dep_time, dep_delay)
#> # A tibble: 336,776 × 19
#>   year month   day dep_time sched_dep_time dep_delay arr_time
#>   <int> <int> <int>   <int>          <int>      <dbl>      <int>
#> 1  2013     7    27        NA            106        NA        NA
#> 2  2013     5     8        445           500       -15        620
#> 3  2013     5     5        446           500       -14        636
#> 4  2013     9     4        446           500       -14        618
#> # ... with 336,772 more rows, and 12 more variables: sched_arr_time <int>,
#> #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
#> #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
#> #   minute <dbl>, time_hour <dtm>

```

The default order is from smallest to largest, numeric to character, and a to z. To reverse this order, use `desc()` to order a column in descending order:

```

arrange(flights, desc(sched_dep_time))
#> # A tibble: 336,776 × 19
#>   year month   day dep_time sched_dep_time dep_delay arr_time
#>   <int> <int> <int>   <int>          <int>      <dbl>      <int>
#> 1  2013     1     1    2353           2359        -6        425
#> 2  2013     1     1    2353           2359        -6        418
#> 3  2013     1     1    2356           2359        -3        425
#> 4  2013     1     2     42           2359        43        518
#> # ... with 336,772 more rows, and 12 more variables: sched_arr_time <int>,
#> #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
#> #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
#> #   minute <dbl>, time_hour <dtm>
arrange(flights, desc(sched_dep_time), dep_delay)
#> # A tibble: 336,776 × 19
#>   year month   day dep_time sched_dep_time dep_delay arr_time
#>   <int> <int> <int>   <int>          <int>      <dbl>      <int>
#> 1  2013    10     2    2341           2359       -18        324
#> 2  2013     9    23    2342           2359       -17        331
#> 3  2013    10    22    2343           2359       -16        347
#> 4  2013     3     4    2343           2359       -16        418
#> # ... with 336,772 more rows, and 12 more variables: sched_arr_time <int>,
#> #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
#> #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
#> #   minute <dbl>, time_hour <dtm>

```

`arrange()` is a straightforward wrapper around `order()` that requires less typing. The previous code is equivalent to:

```
flights[order(flights$sched_dep_time), ]
flights[order(flights$sched_dep_time, decreasing = TRUE, flights$dep_delay), ]
```

Select columns with `select()`

Datasets generally contain numerous columns, but oftentimes you are only interested in a few for a given analysis. `select()` allows you to focus on a useful subset of your data while dropping un-needed columns:

```
# Select columns by name
select(flights, month, day, carrier)
#> # A tibble: 336,776 × 3
#>   month   day carrier
#>   <int> <int>   <chr>
#> 1     1     1     UA
#> 2     1     1     UA
#> 3     1     1     AA
#> 4     1     1     B6
#> # ... with 336,772 more rows
# Select all columns between day and arr_time (inclusive)
select(flights, day:arr_time)
#> # A tibble: 336,776 × 5
#>   day dep_time sched_dep_time dep_delay arr_time
#>   <int>   <int>         <int>     <dbl>   <int>
#> 1     1     517           515         2     830
#> 2     1     533           529         4     850
#> 3     1     542           540         2     923
#> 4     1     544           545        -1    1004
#> # ... with 336,772 more rows
# Select all columns except those from year to day (inclusive)
select(flights, -(year:day))
#> # A tibble: 336,776 × 16
#>   dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay
#>   <int>         <int>         <dbl>   <int>         <int>         <dbl>
#> 1     517           515         2     830           819          11
#> 2     533           529         4     850           830          20
#> 3     542           540         2     923           850          33
#> 4     544           545        -1    1004          1022         -18
#> # ... with 336,772 more rows, and 10 more variables: carrier <chr>,
#> #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>,
#> #   distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

This function works similarly to the `select` argument in `base::subset()`. The dplyr philosophy is to have small functions that do one thing well, so it's its own function.

There are a number of helper functions you can use within `select()`, like `starts_with()`, `ends_with()`, `matches()` and `contains()`. These let you quickly match larger blocks of variables that meet some criterion. See `?select` for more details.

Add new columns with `mutate()`

Data analysis often requires the creation of new variable columns based on values within your dataset. The `mutate()` function allows you to do this:

```
mutate(flights,
  gain = arr_delay - dep_delay,
  speed = distance / air_time * 60)
#> # A tibble: 336,776 × 21
#>   year month   day dep_time sched_dep_time dep_delay arr_time
#>   <int> <int> <int>   <int>         <int>      <dbl>   <int>
#> 1  2013     1     1     517             515         2     830
#> 2  2013     1     1     533             529         4     850
#> 3  2013     1     1     542             540         2     923
#> 4  2013     1     1     544             545        -1    1004
#> # ... with 336,772 more rows, and 14 more variables: sched_arr_time <int>,
#> #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
#> #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
#> #   minute <dbl>, time_hour <dtm>, gain <dbl>, speed <dbl>
```

`dplyr::mutate()` works similarly to `base::transform()`. The key difference between `mutate()` and `transform()` is that `mutate` allows you to refer to columns that you've just created:

```
mutate(flights,
  gain = arr_delay - dep_delay,
  gain_per_hour = gain / (air_time / 60)
)
#> # A tibble: 336,776 × 21
#>   year month   day dep_time sched_dep_time dep_delay arr_time
#>   <int> <int> <int>   <int>         <int>      <dbl>   <int>
#> 1  2013     1     1     517             515         2     830
#> 2  2013     1     1     533             529         4     850
#> 3  2013     1     1     542             540         2     923
#> 4  2013     1     1     544             545        -1    1004
#> # ... with 336,772 more rows, and 14 more variables: sched_arr_time <int>,
#> #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
#> #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
#> #   minute <dbl>, time_hour <dtm>, gain <dbl>, gain_per_hour <dbl>
```

```
transform(flights,
  gain = arr_delay - delay,
  gain_per_hour = gain / (air_time / 60)
)
#> Error: object 'gain' not found
```

Summarise values with `summarise()`

The last verb is `summarise()`. This collapses a dataframe to a single value, based on a function:

```
# mean departure delay
summarise(flights,
  delay = mean(air_time, na.rm = TRUE))
#> # A tibble: 1 × 1
#>   delay
#>   <dbl>
#> 1 150.6865
# shortest flight distance
summarise(flights, min.dist = min(distance))
#> # A tibble: 1 × 1
```

```
#> min.dist
#>      <dbl>
#> 1      17
```

Notice the argument `na.rm = TRUE` within the `summarize` function. If you have missing or NA values within your data, it will cause the summary functions to return NA. Additionally, `summarize` only works with functions that return a single value:

```
summarise(flights, delay.range = range(dep_delay, na.rm = TRUE))
#> Error: expecting result of length one, got : 2
```

However, we can get around this by calling two functions within `summarise`:

```
summarise(flights, min.delay = min(dep_delay, na.rm = TRUE),
          max.delay = max(dep_delay, na.rm = TRUE))
#> # A tibble: 1 × 2
#>   min.delay max.delay
#>   <dbl>      <dbl>
#> 1    -43      1301
```


We will return to the `summarize()` verb below.

Grouped operations

These verbs are useful on their own, but when used in conjunction with the `group_by()` function, the awesomeness of dplyr starts to shine through. It organizes a dataset into specified groups of rows. Verbs are then applied group-by-group within the dataset. Conveniently, this is accomplished using the exact same syntax as above.

In the following example, the entire dataset is grouped into carrier, and then individual planes. A new column for air speed is added, and it is then summarized by counting the number of flights (`count = n()`) and computing the average speed (`mean.speed = mean(speed, na.rm = TRUE)`) and arrival delay (`delay = mean(arr_delay, na.rm = TRUE)`). The results are then displayed using `ggplot`.

```
#filter out canceled flights
not_cancelled <- filter(flights, !is.na(dep_time))
by_carrier <- group_by(not_cancelled, carrier, tailnum)
air_speed <- mutate(by_carrier, speed = distance / air_time * 60)
delay <- summarise(air_speed,
  count = n(),
  distance = mean(distance, na.rm = TRUE),
  speed = mean(speed, na.rm = TRUE))
delay <- filter(delay, count > 30)
ggplot(delay, aes(distance, speed)) +
  geom_point(aes(color = carrier, size = count), alpha = 0.5)
```



prisingly, flights that travel a shorter distance do not reach their maximal cruising speed.

Chaining

The dplyr function calls don't have any side-effects (unlike some base functions), making it easy to explore your data in an interactive way. However, one disadvantage of this is it doesn't lead to very succinct code, particularly if you want to perform many operations at once. You can do it step-by-step, saving a new object each time:

```
# calculate flights either whose arrival or departure were delayed > 30 minutes
a1 <- group_by(flights, year, month, day)
a2 <- select(a1, arr_delay, dep_delay)
a3 <- summarise(a2,
  arr = mean(arr_delay, na.rm = TRUE),
  dep = mean(dep_delay, na.rm = TRUE))
a4 <- filter(a3, arr > 30 | dep > 30)
```

However this can lead to many problems. Giving objects appropriate names can be difficult (e.g. object names in ggplot example above). Or naming them chronologically (as above), it can be difficult to remember which object is which (was the summarized object a2 or a3?...). Especially if you want to do the same thing over and over on different data sets or subsets of observations.

Alternatively, if you don't want to save the intermediate results, you need to wrap the function calls inside each other:

```
# calculate flights either whose arrival or departure were delayed > 30 minutes
filter(
  summarise(
```



```

select(
  group_by(flights, year, month, day),
  arr_delay, dep_delay
),
arr = mean(arr_delay, na.rm = TRUE),
dep = mean(dep_delay, na.rm = TRUE)
),
arr > 30 | dep > 30
)
#> Adding missing grouping variables: `year`, `month`, `day`
#> Source: local data frame [49 x 5]
#> Groups: year, month [11]
#>
#>   year month   day      arr      dep
#>   <int> <int> <int>   <dbl>   <dbl>
#> 1  2013     1    16 34.24736 24.61287
#> 2  2013     1    31 32.60285 28.65836
#> 3  2013     2    11 36.29009 39.07360
#> 4  2013     2    27 31.25249 37.76327
#> # ... with 45 more rows

```

This is difficult to write and to read, and not intuitive because the order of the operations is from inside to out. The arguments are a long way away from the function. for example the arguments for the first `filter()` function above are the last thing within the full call (`arr > 30 | dep > 30`).

In order to get around this and write intuitive, easy to read code, we need to chain the function calls together. But before we do this, we need to introduce the pipe operator `%>%`

The pipe operator: `%>%`

The pipe operator is from the `magrittr` package, but is included in the `dplyr` package, so no need to load this library separately.

The pipe operator turns `x %>% f(y)` into `f(x,y)`.

I think that this is a bit confusing to think about, but when you see some examples the power and ease of use becomes obvious.

```

# a silly example
x <- seq(10)
max(x)
#> [1] 10
# is the same as
x %>% max
#> [1] 10

mean(x)
#> [1] 5.5
x %>% mean
#> [1] 5.5

```

You can see that both syntaxes give the same result. Now, let's go back our example looking at arrival and departure delays > 30 minutes

```

# original method, saving new object at each step
a1 <- group_by(flights, year, month, day)

```

```

a2 <- select(a1, arr_delay, dep_delay)
a3 <- summarise(a2,
  arr = mean(arr_delay, na.rm = TRUE),
  dep = mean(dep_delay, na.rm = TRUE))
a4 <- filter(a3, arr > 30 | dep > 30)

# now using the %>% operator
a5 <- flights %>% group_by(year, month, day) %>%
  select(arr_delay, dep_delay) %>%
  summarise(
    arr = mean(arr_delay, na.rm = TRUE),
    dep = mean(dep_delay, na.rm = TRUE)) %>%
  filter(arr > 30 | dep > 30)

identical(a4, a5)
#> TRUE

```

It is helpful to say “then” when you see a %>% operator. e.g. take object “flights”, *then* group by year, month and day, *then* select variables arr_delay and dep_delay *then* ...

The pipe operator in combination with dplyr allows you to quickly examine your data and explore interesting results. One of my favorite aspects of this is you can answer interesting questions almost immediately.

“I wonder what the smallest arrival delay on Christmas day was, and how far the plane travelled?”

```

flights %>%
  filter(month == 12, day == 25) %>%
  select(dep_delay, distance) %>%
  arrange(dep_delay)
#> # A tibble: 719 × 2
#>   dep_delay distance
#>   <dbl>      <dbl>
#> 1      -23       762
#> 2      -16      1389
#> 3      -15      2402
#> 4      -15      1608
#> # ... with 715 more rows
# note that negative values indicate an early arrival

```

If you want to also know the destination:

```

flights %>%
  filter(month == 12, day == 25) %>%
  select(dep_delay, distance, dest) %>%
  arrange(dep_delay)
#> # A tibble: 719 × 3
#>   dep_delay distance dest
#>   <dbl>      <dbl> <chr>
#> 1      -23       762 ATL
#> 2      -16      1389 DFW
#> 3      -15      2402 SEA
#> 4      -15      1608 SJU
#> # ... with 715 more rows

```

What carrier has the longest flights on average?

```

flights %>%
  group_by(carrier) %>%
  summarise(mean.dist = mean(distance),
            count = n()) %>%
  arrange(desc(mean.dist))
#> # A tibble: 16 × 3
#>   carrier mean.dist count
#>   <chr>     <dbl> <int>
#> 1     HA  4983.000   342
#> 2     VX  2499.482  5162
#> 3     AS  2402.000   714
#> 4     F9  1620.000   685
#> # ... with 12 more rows

```

What questions do you want to answer? Think about what you can ask with this dataset, and then go step-by-step to answer the question.

You can also pipe objects into other functions outside the dplyr package:


```

flights %>%
  filter(carrier == "HA") %>%
  dim()
#> [1] 342  19

# pipe results into write.csv()
flights %>%
  group_by(carrier, dest) %>%
  select(dest, carrier, dep_delay) %>%
  summarize(mean.delay = mean(dep_delay, na.rm = T),
            count = n()) %>%
  arrange(carrier, desc(mean.delay)) %>%
  write.csv("mean delay by carrier and dest.csv", row.names = F)

# revisit our ggplot example above
# pipe results into ggplot()
flights %>%
  filter(!is.na(dep_time)) %>%
  group_by(carrier, tailnum) %>%
  mutate(speed = distance / air_time * 60) %>%
  summarize(
    count = n(),
    distance = mean(distance, na.rm = T),
    speed = mean(speed, na.rm = T)) %>%
  filter(count > 30) %>%
  ggplot(aes(distance, speed)) +
  geom_point(aes(color = carrier, size = count), alpha = 0.5)

```



Quick note, if you don't know where to find the .csv file you created in the example above, you can check your working directory with `getwd()`. This is also a good time to mention that if you're not already using projects in Rstudio, I highly recommend it. See Chapter 6 "Workflow: Projects" in *R for Data Science* for more information. <http://r4ds.had.co.nz/workflow-projects.html/>

Two table verbs

Oftentimes the data you need is not in only one dataset. For example, the above results showing which destinations have the largest average departure delay may be interesting, but who has all the carrier and airport abbreviations memorized?

The above result would be more useful to lay persons if it had more specific information. For this we can use the two table verbs of `dplyr`.

The `airports` and `airlines` datasets within the `nycflights13` data has some useful information.

```
#> # A tibble: 6 × 2
#>   carrier      name
#>   <chr>      <chr>
#> 1     9E Endeavor Air Inc.
#> 2     AA American Airlines Inc.
#> 3     AS  Alaska Airlines Inc.
#> 4     B6   JetBlue Airways
#> # ... with 2 more rows
#> # A tibble: 6 × 8
#>   faa      name    lat    lon  alt  tz  dst
#>   <chr>   <chr>   <dbl> <dbl> <int> <dbl> <chr>
```

```
#> 1 04G Lansdowne Airport 41.13047 -80.61958 1044 -5 A
#> 2 06A Moton Field Municipal Airport 32.46057 -85.68003 264 -6 A
#> 3 06C Schaumburg Regional 41.98934 -88.10124 801 -6 A
#> 4 06N Randall Airport 41.43191 -74.39156 523 -5 A
#> # ... with 2 more rows, and 1 more variables: tzone <chr>
```

There are a number of joining commands within `dplyr`, primarily differing in how they treat missing data.

- Join rows from `b` matching to `a`, keeping all rows from `a`: `left_join(a, b, by = "x1")`
- Join rows from `a` matching to `b`, keeping all rows from `b`: `right_join(a, b, by = "x1")`
- Join data, only keeping rows in `a` and `b`: `inner_join(a, b, by = "x1")`
- Join data, keep all rows in both datasets: `full_join(a, b, by = "x1")`

The Rstudio data wrangling cheat sheet has very useful diagrams explaining this visually.

All of these can be accomplished with using arguments within the `base::merge()` function, but as noted previously, one of the core features of `dplyr` is to have “small” functions that do one thing well.

The most commonly used joining verb is `left_join` because it retains all rows in the `a` dataset, whether they have a match in `b` or not.

```
flights2 <- left_join(flights, airlines, by = "carrier") %>% rename(airline = name)
head(flights2)
#> # A tibble: 6 × 20
#>   year month   day dep_time sched_dep_time dep_delay arr_time
#>   <int> <int> <int>   <int>         <int>       <dbl>   <int>
#> 1  2013     1     1     517             515         2     830
#> 2  2013     1     1     533             529         4     850
#> 3  2013     1     1     542             540         2     923
#> 4  2013     1     1     544             545        -1    1004
#> # ... with 2 more rows, and 13 more variables: sched_arr_time <int>,
#> #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
#> #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
#> #   minute <dbl>, time_hour <dtm>, airline <chr>
```

We used the function `dplyr::rename()` to rename the `name` column to `airline`. This is so when we add the airport names there are not multiple columns with the same id.

The `airports` dataset has a lot of information, but we’re really only interested in the airport names. We can combine single table and two table verbs to only join the airport name with the airport code.

```
flights3 <- airports %>%
  # select only columns of interest
  select(faa, name) %>%
  # join the two datasets
  left_join(flights2, ., by = c("dest" = "faa")) %>%
  rename(dest.name = name)
```

You may have noticed something tricky in that last section of code. What does the `.`, represent? So far, all of our piping operations have used the output of one function as the first argument in the next function. But in the above example, we want the output of the first function as the *second* argument in the second function. In order to do this we use the `.` to let R know that that’s where we want the output to go. The reason we want it as the second argument is that `left_join()` keeps all of the *first* objects rows, and gets rid of any non-matched in the second dataset. An alternative to using the `.` notation would have been to use `right_join()`:

```
flights3.b <- airports %>%
  select(faa, name) %>%
```

```
right_join(flights2, by = c("faa" = "dest")) %>%
  rename(dest.name = name)
```

Notice that the order of “faa” and “dest” has been switched. If you tried it with the original order of “dest” = “faa” it would have thrown an error because R would be looking for a column named “dest” in the first dataset which doesn’t exist.

Also important to note that we need to be careful here, as the `right_join()` default keeps all of the `b` dataset and drops any non-matches in `a`.

All of this is to say that the order of data inputs affects the order of your code. The syntaxes are all so similar that it’s very easy to write code that reads well and looks like it should work, but can throw unexpected errors. But back to the task at hand; finishing adding airport names to the `flights3` dataset.

```
flights4 <- airports %>%
  select(faa, name) %>%
  left_join(flights3, ., by = c("origin" = "faa")) %>%
  rename(origin.name = name)
```

Now we can re-run our example above and have more useful information. Make sure to change the variable names so you’re working with our updated objects and variables.

```
flights4 %>%
  group_by(airline, dest.name) %>%
  select(dest.name, airline, dep_delay) %>%
  summarize(mean_delay = mean(dep_delay, na.rm = T),
  count = n()) %>%
  arrange(airline, desc(mean_delay))
#> Source: local data frame [308 x 4]
#> Groups: airline [16]
#>
#>           airline                dest.name mean_delay
#>           <chr>                  <chr>         <dbl>
#> 1 AirTran Airways Corporation Akron Canton Regional Airport 20.820878
#> 2 AirTran Airways Corporation Hartsfield Jackson Atlanta Intl 18.445609
#> 3 AirTran Airways Corporation General Mitchell Intl -1.709091
#> 4 Alaska Airlines Inc. Seattle Tacoma Intl 5.804775
#> # ... with 304 more rows, and 1 more variables: count <int>
```

Plug for “tidy” data

“Tidy datasets are all alike, but every messy dataset is unique in its own way” -*Hadley Wickham*

You may have noticed that manipulating the `flights` dataset was extremely easy and intuitive using the `dplyr` package. One of the reasons for this, is that `dplyr` is optimized to work with “tidy” data, and the `flights` data is already in the tidy format.

So what is tidy data? For an in depth look, see the *Tidy Data* paper <http://www.jstatsoft.org/v59/i10/paper/>.

Briefly, a tidy dataset is where all columns are a single variable, each observation is its own row, and each value has its own cell.

However, many datasets are messy, and decidedly not tidy. This may be due to how data was collected, the output readings of our instruments, the fact that it’s easier to enter data in a certain way, etc.

If possible, it is easiest to enter all of your data in a tidy format, foregoing any need to manipulate our data in R, or, god forbid, in Excel. However, if you’re working with old datasets, or have to enter data in an

un-tidy way, have no fear, as there's a package for that.

The `tidyr` package was designed to turn messy data into tidy data, with surprisingly few commands (see the data wrangling cheat sheet we've been talking about). Unfortunately we don't have time to go into this now but there are a number of resources available for further reading and example datasets available online (see resources below).

One more quick plug for data formatting and manipulation in R, as opposed to Excel. One of the best utilities of R (or a similar language) is the ability to make your analysis completely reproducible. It may be easier now to just do your data formatting in excel, but it is fraught with dangers. It's easy to change an equation in a cell, or reference the wrong cell, or accidentally copy a cell's equation instead of the value etc., and working backwards to find the error can be a literal nightmare. Also, changing your data's format in excel can be dangerous. You may have done it right the first time, but there's no paper trail, no breadcrumbs to follow backwards to find your error. Likewise, if you return to your data much later, it is amazingly easy to forget how exactly you moved things around, or what values were calculated how.

However, doing all of this in R allows you to leave clear signposts for exactly what you did and the order you did it in. This also allows you to leave your raw data alone. I can't tell you how many times I've messed something up and had to start over with the raw data. But creating a data formatting script allows me to find where the mistake is, fix it, and then reimport the raw data, source the entire document `Ctrl + Shift + S` and be on my merry way. This is also useful if you have numerous excel sheets that all need the same functions called on them (e.g. HOB0 data logger output files).

This concludes the introduction to the `dplyr` package. Hopefully it has given you a taste of the power and ease of using `dplyr` with tidy data, and provided you with enough resources to continue on and learn more if you need it.

Resources

`dplyr` now has its own website! Work in progress as of 5 Apr 2017 <http://dplyr.tidyverse.org/>

`dplyr` vignettes `browseVignettes("dplyr")`

"R for Data Science by Hadley Wickham & Garrett Grolemund (O'reilly). Copyright 2017. 978-1-491-91039-9." Covers `dplyr`, `tidyr`, as well as the full tidyverse. Also a great general introduction to R, data analysis, workflow, programming, etc. Available free online at <http://r4ds.had.co.nz/>

`dplyr` introduction by Sean Anderson Covers many of the ideas introduced here, but using a large 'mammals.txt' dataset. <https://www.dropbox.com/sh/kn6fwme5e9np8l/AAAv2KMlapwAs6Yd30UbYV9ea?dl=0>

Rstudio data wrangling cheat sheet (`dplyr` & `tidyr`) This is an amazing resource that I have open pretty much everytime I turn R on. <https://www.rstudio.com/wp-content/uploads/2015/02/data-wrangling-cheatsheet.pdf/>.

`tidyr` vignettes `browseVignettes("tidyr")`

`tidyr` introduction by Teja Kodali <https://datascienceplus.com/data-manipulation-with-tidyr/>

Generally good resources <https://blog.rstudio.org/> Aggregates a number of blogs. I follow them on twitter and can't tell you how many times I've randomly come across an article that provides a simple solution to an issue that that I had deemed impossible.

Online courses: Many.

I paid for a few months subscription to DataCamp <https://www.datacamp.com/> and found their courses to be very useful. They have a whole course on `dplyr`, although when I took it it was a bit dated; at the end they were talking about "introducing" some function soon which I had already been using for a while at the time, but still a good place to start.

Twitter: @Rbloggers #rstats

Google. Seriously, if you have a question about *ANYTHING* in R, just G.T.S.
justin.pomeranz@pg.canterbury.ac.nz twitter @NotAPomegranite